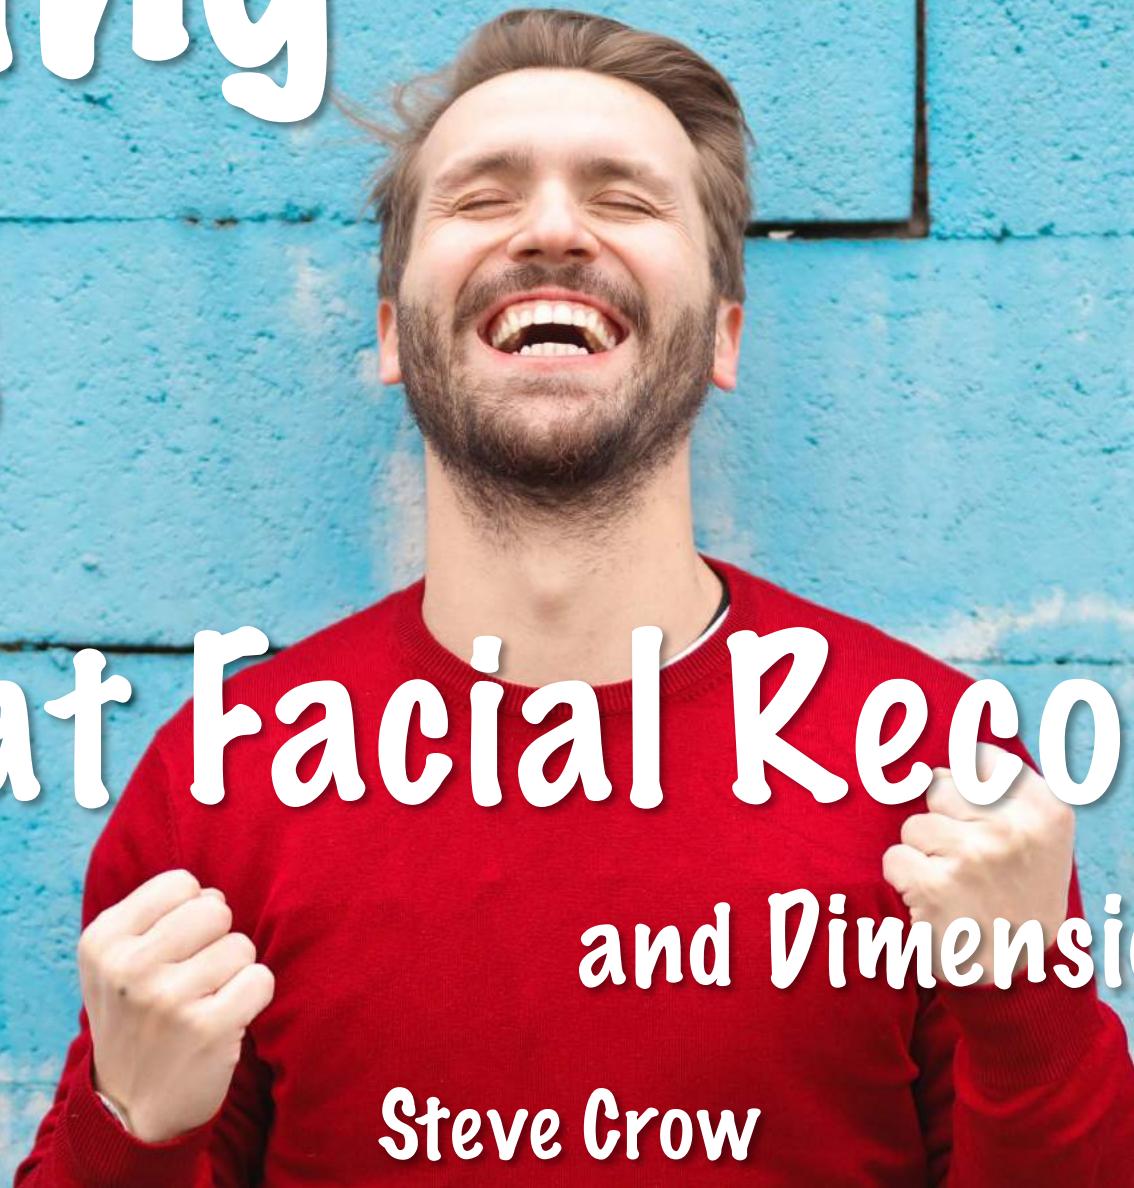


# Making Faces

A Look at Facial Recognition  
and Dimension Reduction

Steve Crow



# Machine Learning

A method of data analysis that automates model building. The goal is to build a mathematical model of sample data from which future predictions can be made.



# Training Set

A set of example data used for learning.

# (Semi) Supervised Learning

Supervised Learning uses a combination of input and output data in the training set to build a mathematical model.

# Unsupervised Learning

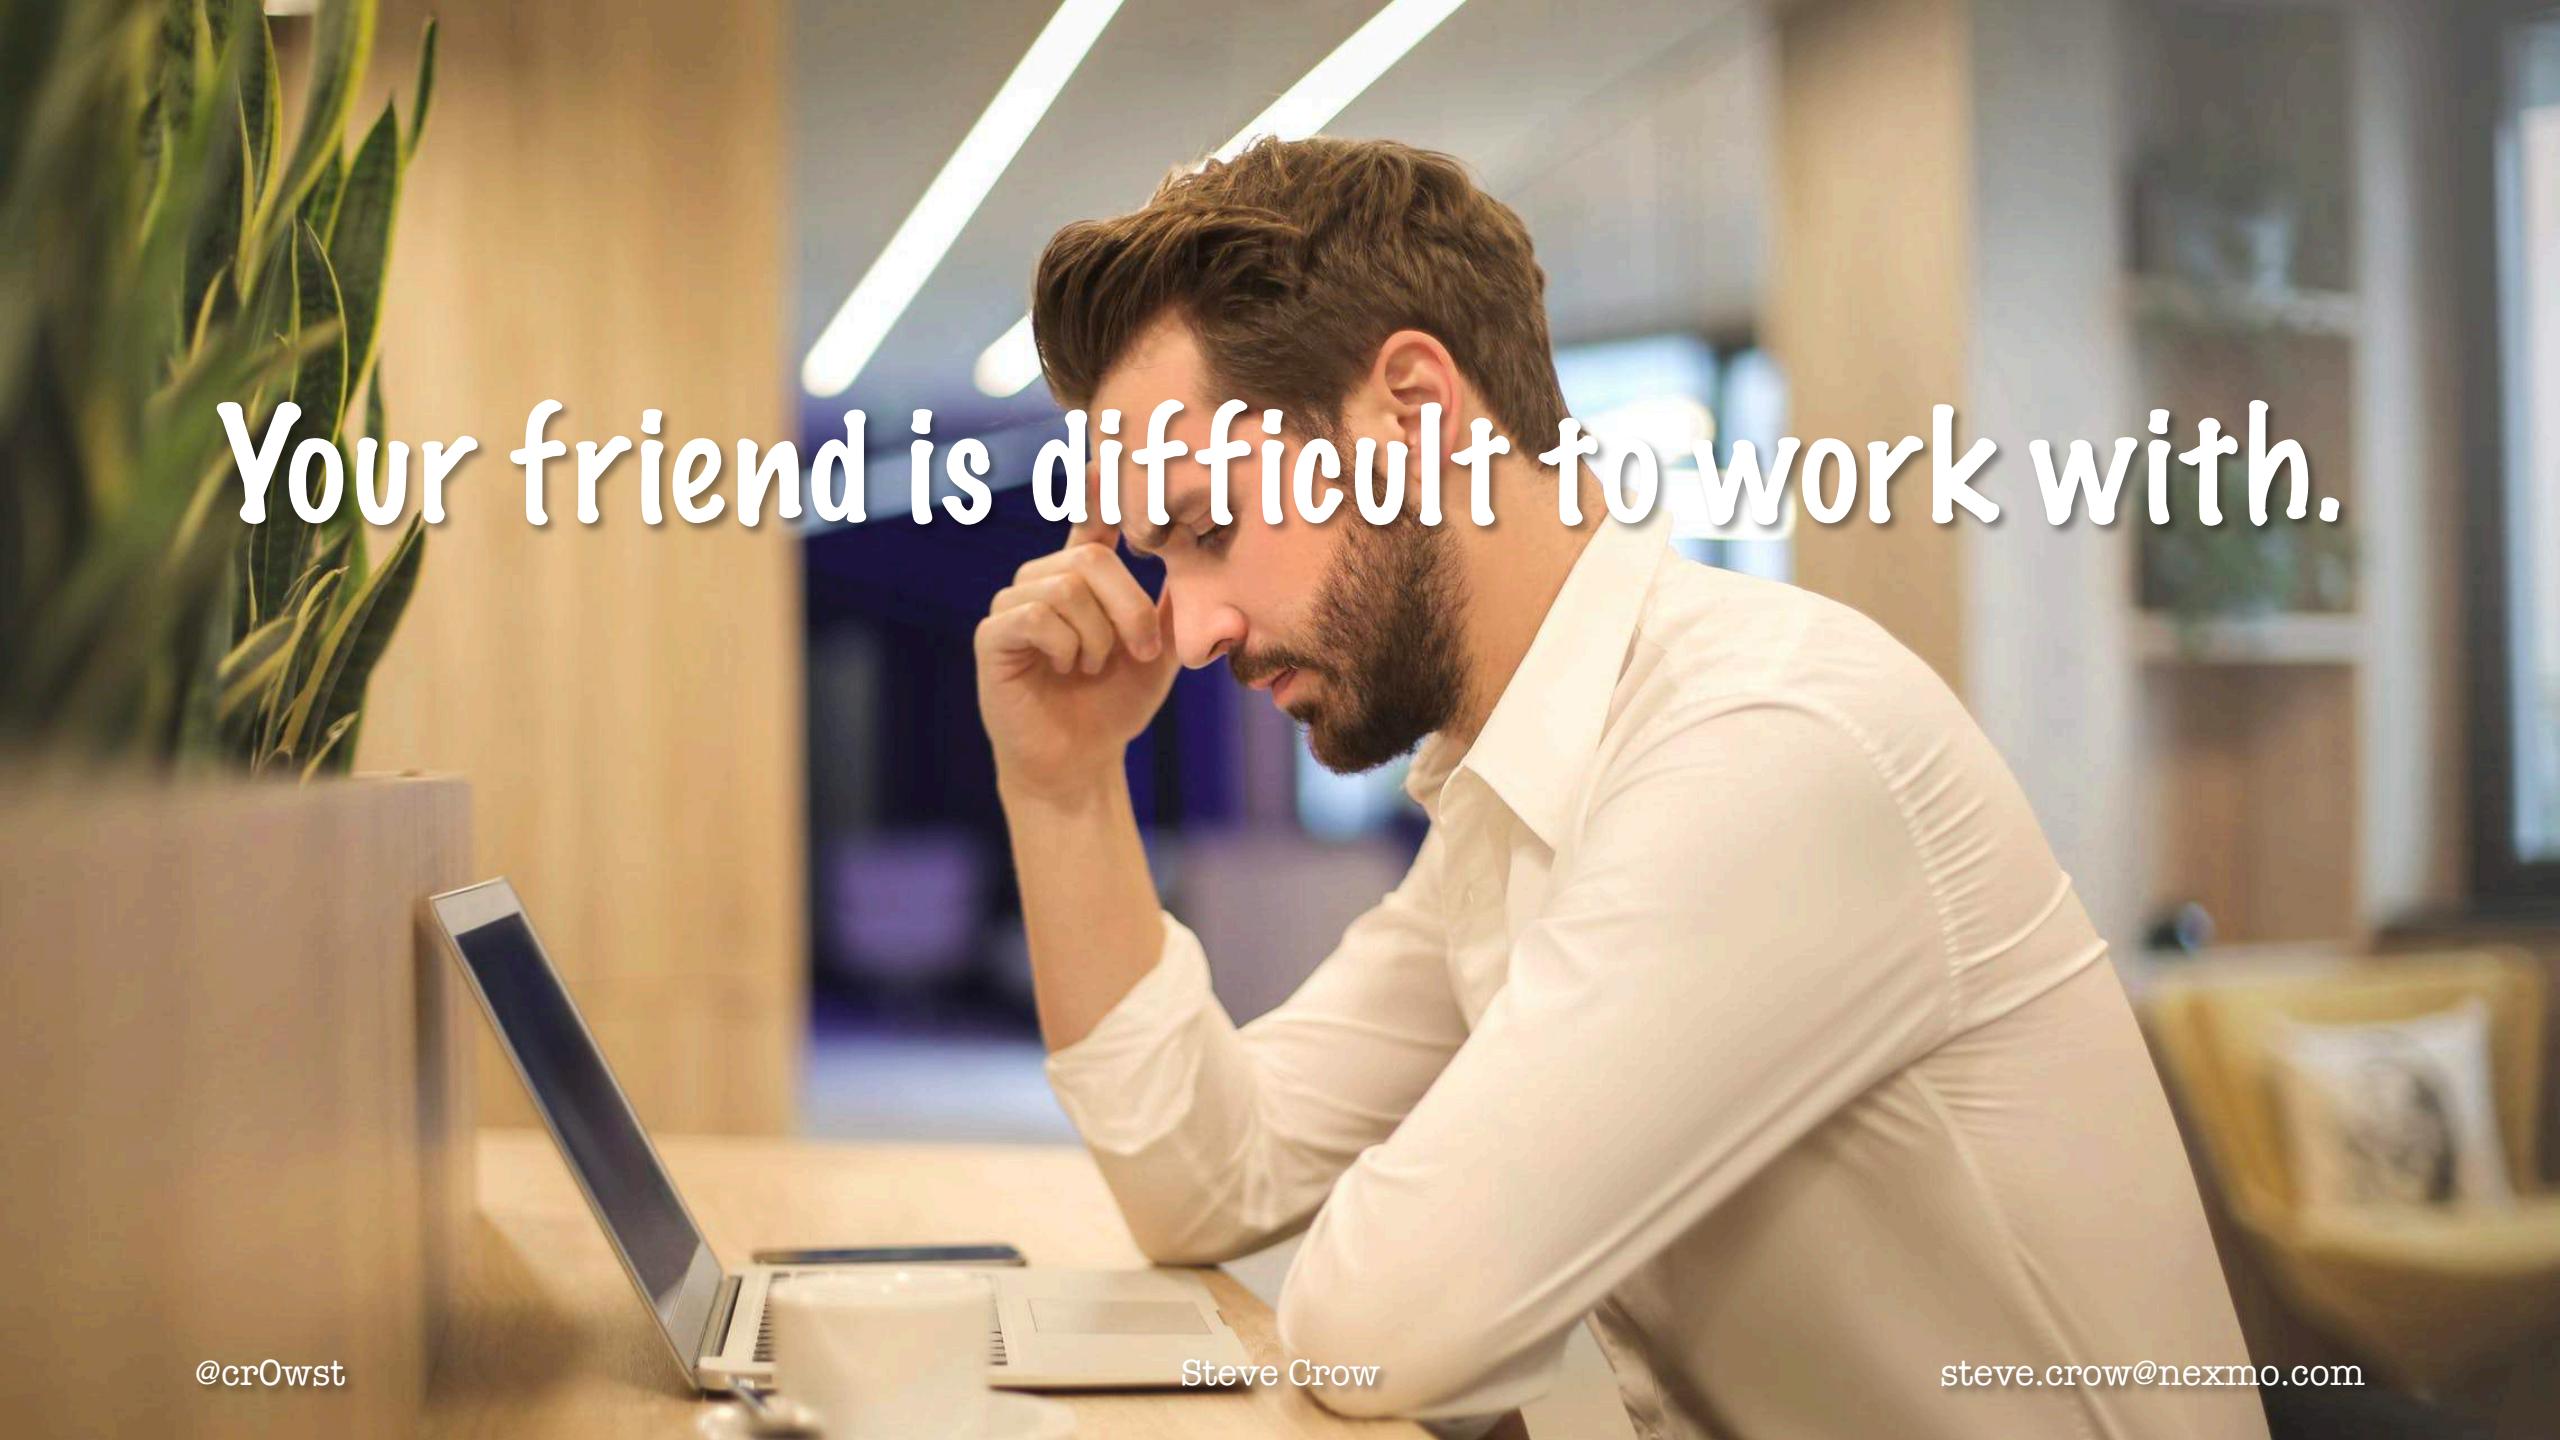
Unsupervised learning uses only inputs in the dataset and finds structures or grouping in the data.



Your friend is throwing a birthday  
party!



# You're in charge of the music!

A photograph of a man with a beard and short brown hair, wearing a white long-sleeved shirt. He is sitting at a light-colored wooden desk, leaning forward with his right hand resting against his forehead and his fingers near his eyes, suggesting stress or exhaustion. In front of him is an open laptop. To his left, a large potted plant with long, thin, variegated leaves sits on a shelf. The background is a blurred office environment with fluorescent lighting.

Your friend is difficult to work with.

A large magnifying glass is positioned in the center of the frame, its circular lens focused on a blurred background of a sunset or sunrise over water. The background has warm orange and yellow tones at the top transitioning to cooler blues and purples at the bottom.

# How do we find music they like?

# Which songs to pick?



# Musical Features

acousticness, danceability, energy, instrumentalness,  
liveness, loudness, speechiness, valence, tempo, time signature,  
key, duration...

# Principal Component Analysis

A process for compressing, extracting features, and reducing dimensions.



# Feature Elimination

*Reduce the feature space by eliminating features.*

An aerial photograph of a vast open-pit coal mine. A massive yellow mining truck is positioned in the center-left, with its long conveyor belt stretching towards the bottom right. The mine's edges are terraced, showing the scale of extraction. The ground is a mix of brown earth and dark, extracted rock.

# Feature Extraction

*Create new features as combination of existing features in an attempt to reduce any sort of redundancy.*

# When to use PCA?

Use Principal Component Analysis when you want to reduce the number of variables, but don't necessarily know which ones to remove completely.

# When to use PCA?

Use Principal Component Analysis when you want to make sure your variables are independent of one another.

# When to use PCA?

Use Principal Component Analysis when you are okay with making your independent variables less interpretable.

# Let's try it out!

# Song Features

Spotify has an API that allows us to access song features. The features that we are going to isolate are energy and loudness.

# Step 1: Create the Training Set

```
> source("/Users/crow.Repositories/making-faces/song example.R")
```

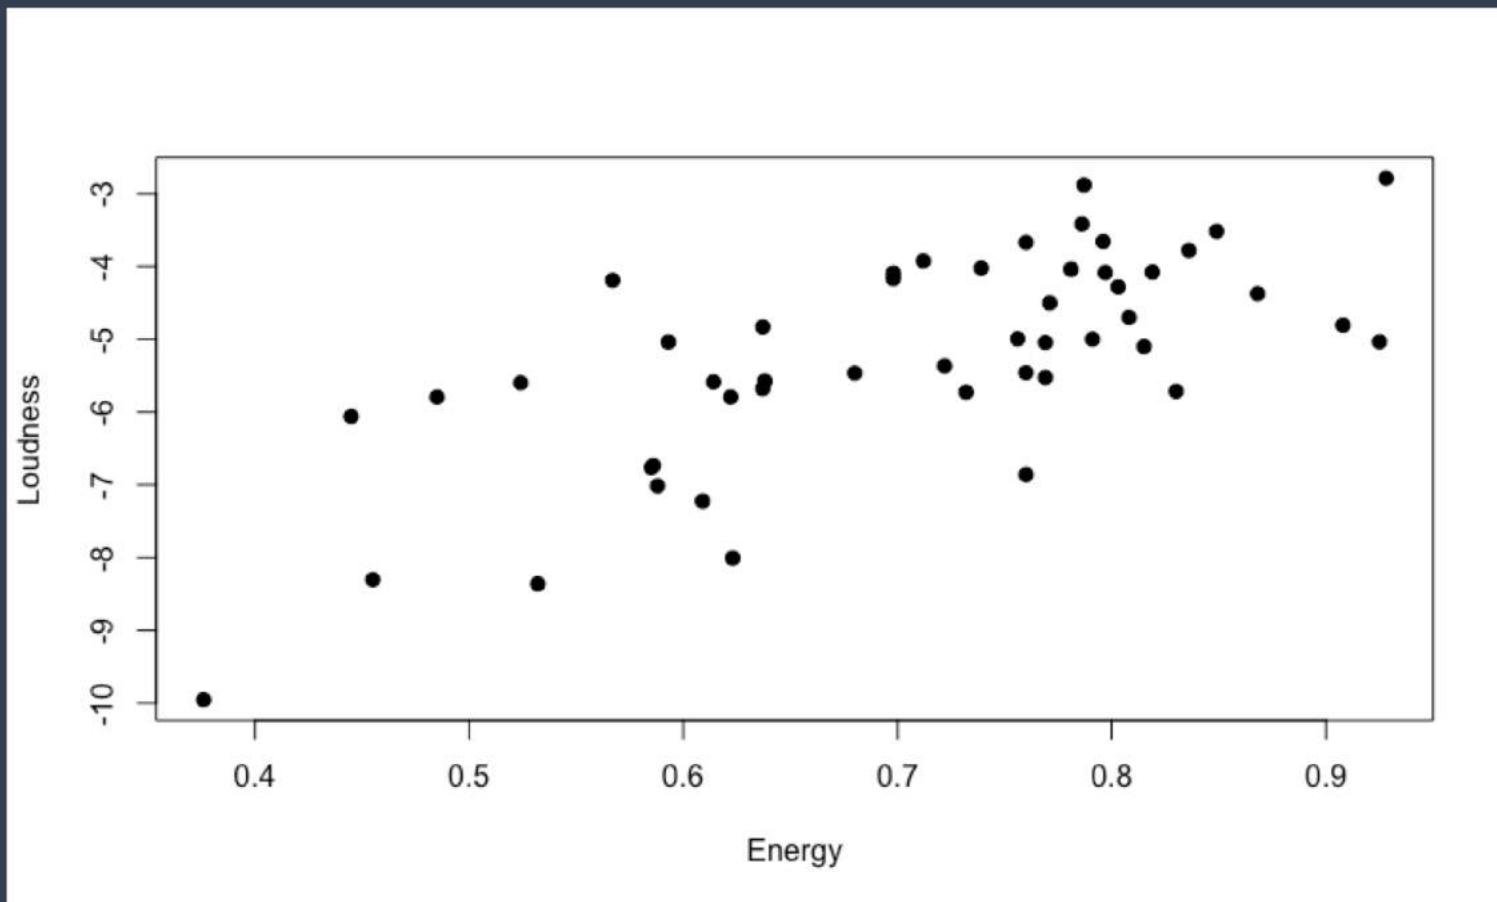
		name	danceability	energy	key	loudness	speechiness	acousticness	instrumentalness	liveness	valence	tempo
1		Love The Way You Lie	0.749	0.925	10	-5.034	0.2270	0.24100	0.00e+00	0.5200	0.641	86.989
2	Hymn For The Weekend - Seeb Remi		0.565	0.849	0	-3.516	0.0517	0.00868	4.68e-06	0.1200	0.427	102.034
3		Uptown Funk	0.856	0.609	0	-7.223	0.0824	0.00801	8.15e-05	0.0344	0.928	114.988
4		We Can't Stop	0.613	0.622	1	-5.794	0.0334	0.00882	0.00e+00	0.3700	0.484	80.003
5	This Is What You Came For		0.630	0.928	9	-2.787	0.0331	0.19900	1.24e-01	0.1480	0.465	123.963
6		One Last Time	0.628	0.593	8	-5.036	0.0323	0.09300	1.65e-06	0.0960	0.104	125.026
7		Chandelier	0.399	0.787	1	-2.880	0.0499	0.01970	6.07e-05	0.0685	0.572	117.089
8		Never Forget You	0.583	0.732	11	-5.728	0.0457	0.00312	9.86e-06	0.2690	0.276	145.992
9		Cool for the Summer	0.584	0.614	5	-5.587	0.0363	0.00506	9.70e-05	0.0802	0.280	114.075
10		Lush Life	0.694	0.712	7	-3.923	0.0460	0.13300	0.00e+00	0.2110	0.799	98.022
11		Thinking Out Loud	0.781	0.445	2	-6.061	0.0295	0.47400	0.00e+00	0.1840	0.591	78.998
12	You Know You Like It		0.426	0.722	5	-5.369	0.1960	0.01130	1.27e-06	0.2620	0.250	196.133
13		Unforgettable	0.726	0.769	6	-5.043	0.1230	0.02930	1.01e-02	0.1040	0.733	97.985
14		Elastic Heart	0.421	0.791	9	-4.998	0.0496	0.01170	1.48e-05	0.1460	0.499	130.075
15	I Really Like You		0.622	0.808	5	-4.701	0.0418	0.00543	5.63e-05	0.2150	0.599	122.121
16		Closer	0.748	0.524	8	-5.599	0.0338	0.41400	0.00e+00	0.1110	0.661	95.010
17		Love Yourself	0.607	0.376	4	-9.954	0.4530	0.85600	0.00e+00	0.2850	0.545	102.541
18	Cake By The Ocean		0.771	0.760	4	-5.459	0.0514	0.14700	0.00e+00	0.0530	0.881	119.000
19		Budapest	0.717	0.455	5	-8.303	0.0276	0.08460	0.00e+00	0.1100	0.389	127.810
20		Sorry	0.654	0.760	0	-3.669	0.0450	0.07970	0.00e+00	0.2990	0.410	99.945
21		Bang Bang	0.706	0.786	0	-3.417	0.0909	0.26000	0.00e+00	0.3800	0.749	150.035
22		Starboy	0.678	0.588	7	-7.015	0.2760	0.14100	6.35e-06	0.1370	0.486	186.005

# Step 1: Create the Training Set



```
songs = read.csv('songdata.csv')[c('energy', 'loudness')]  
training_set = t(data.matrix(songs))
```

# Step 1: Song Space



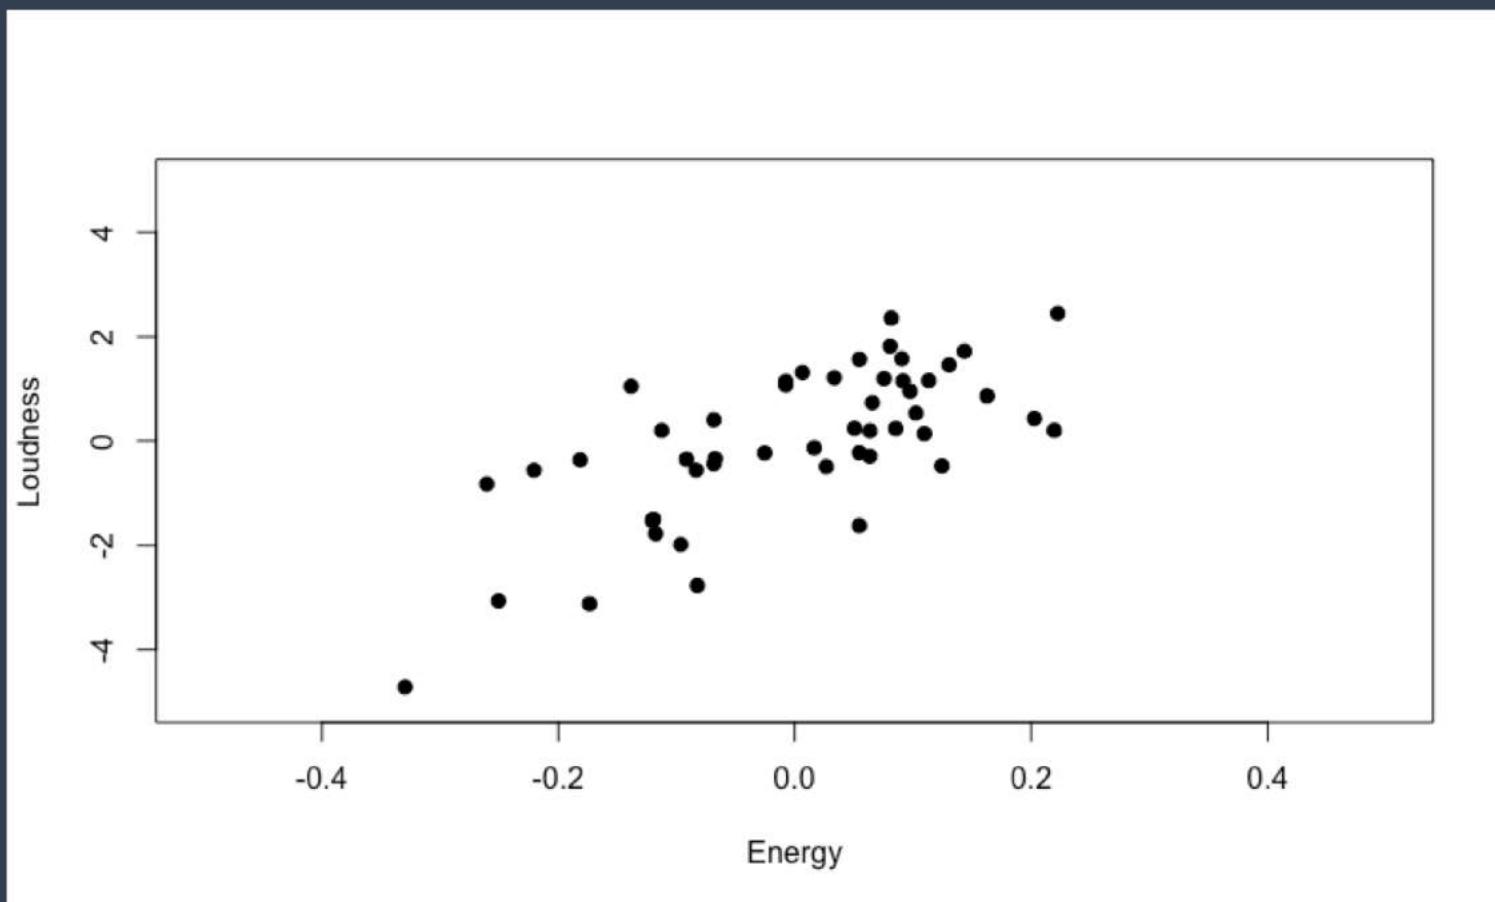
# Step 2: Normalize the Data

- Find the average energy and loudness
- Subtract from each energy and loudness to normalize the data.

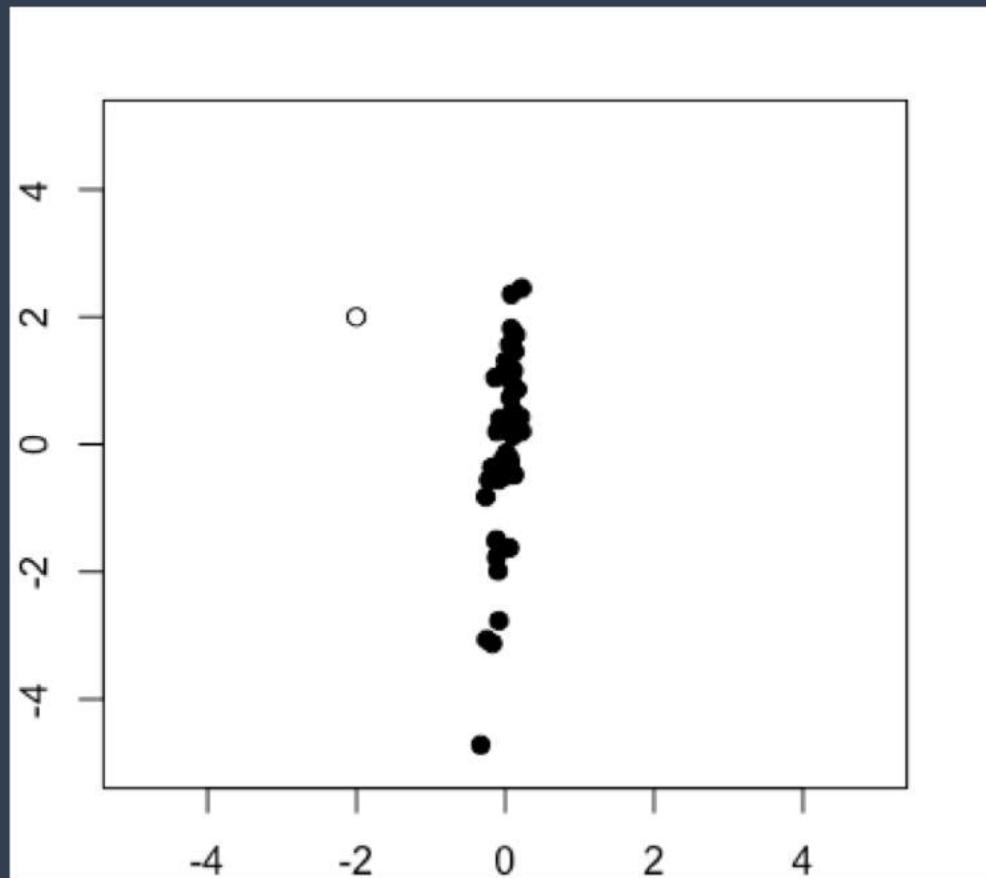


```
centered = training_set - rowMeans(training_set)
```

# Step 2: Normalize the Data

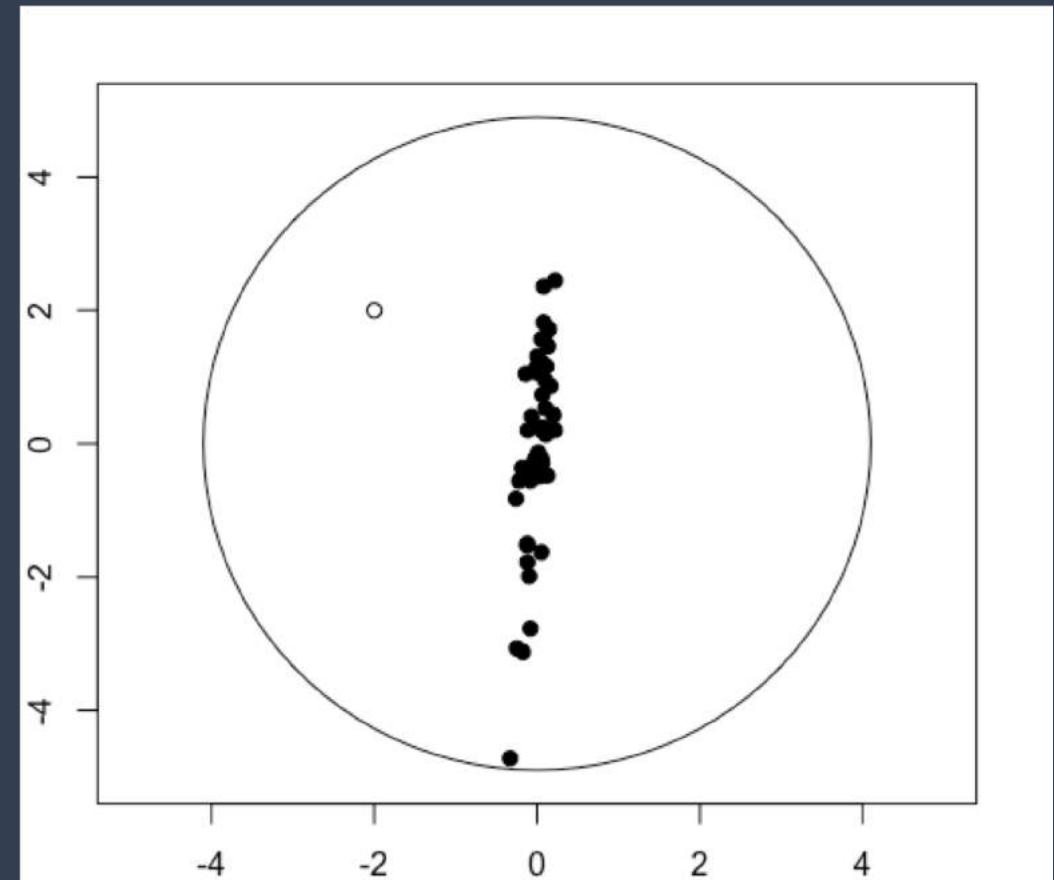


# Does the song belong on the playlist?



# Create a Boundary

- If the point is inside of the circle then the song belongs on the playlist.
- Is there something wrong with this?



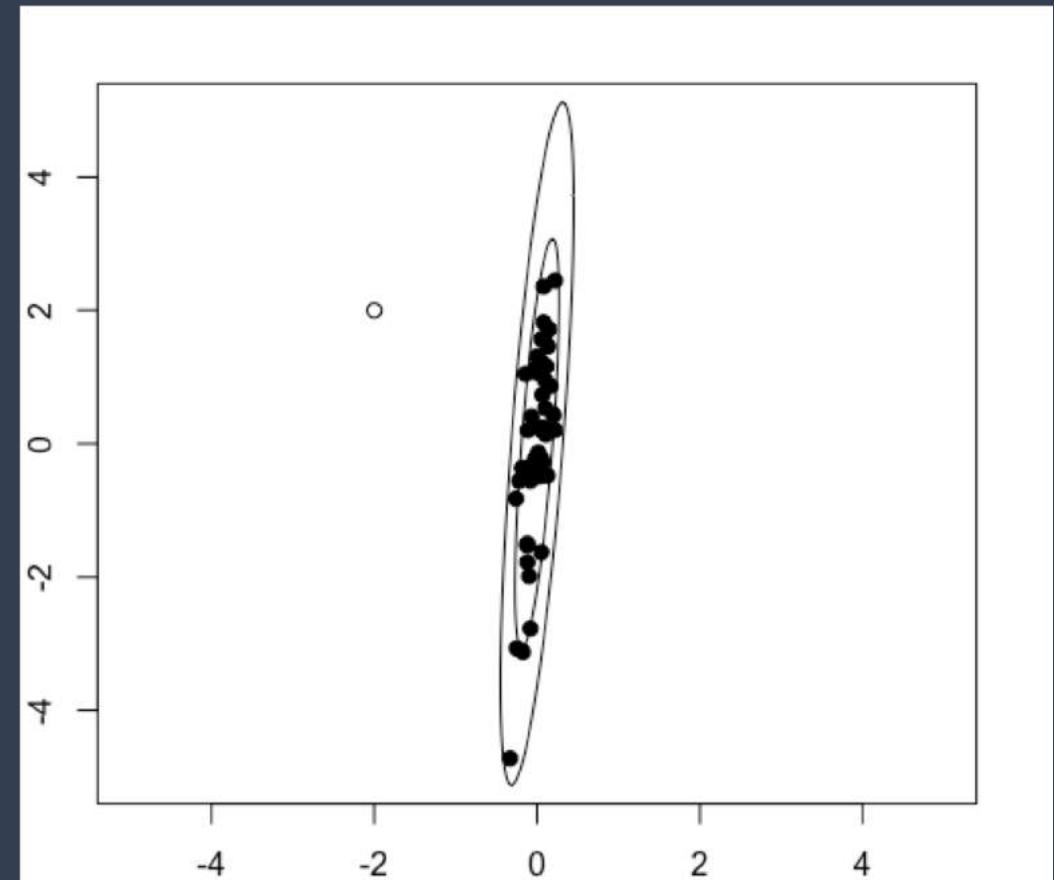
A photograph of a black wrought-iron fence in a park. The fence has vertical bars and decorative curved armrests. In the background, there's a paved path curving to the right, surrounded by green trees and bushes. The lighting suggests it's either early morning or late afternoon.

# We need a better boundary.

Instead of a circle, maybe we can find a better shape to fit the dataset that uses the actual dataset.

# Create a Better Boundary

- If the point is inside of the ellipse then the song belongs on the playlist.
- The axes of the ellipse point to the direction of least and most variance.
- The size of the ellipse depends on how confident we want to be in fitting the data.



# The Singular Value Decomposition

One of the most known and widely used matrix decompositions. All matrices have an SVD and it is commonly used in compression, denoising, and data reduction.

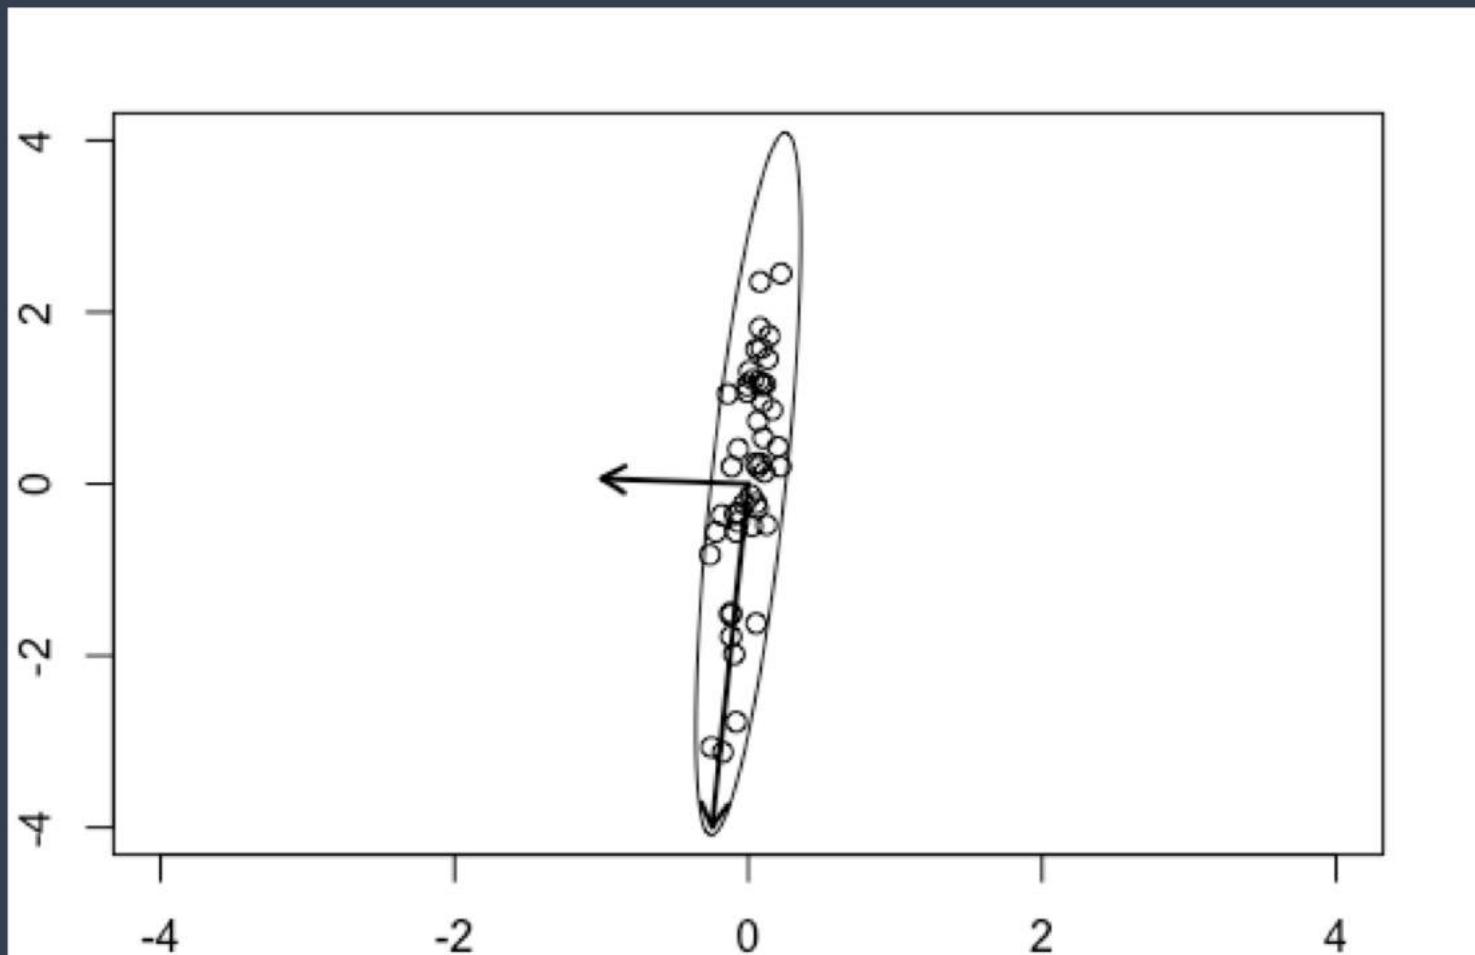
# The Singular Value Decomposition

$$A = U\Sigma V^T$$

# Step 3: The Singular Value Decomposition

- Let  $M$  equal our Song Space Matrix and  $n$  the number of songs.
- Then  $\text{svd} \left( \frac{M}{\sqrt{n}} \right) = U\Sigma V^T$
- The columns of  $U$  are vectors which point in the direction of greatest to least variance.
- The diagonal values of  $\Sigma$  give insight into how significant each of the principal components are.

# Step 3: The Singular Value Decomposition



# Step 3: The Singular Value Decomposition



```
svd = svd(centered / sqrt(ncol(training_set))))
```

# The Singular Value Decomposition

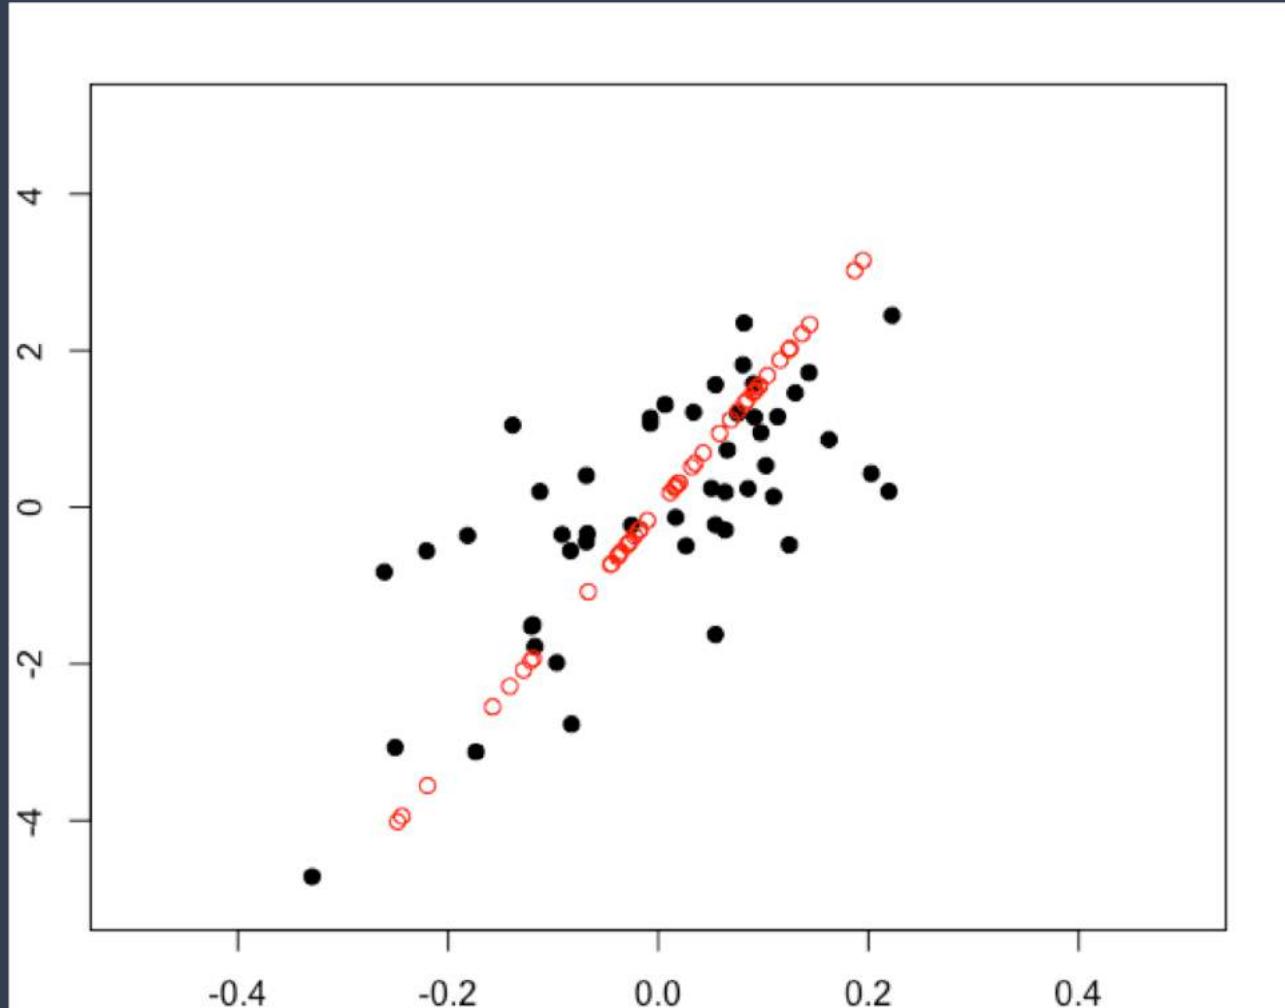
$$A = U\Sigma V^T$$

# Matrix Multiplication

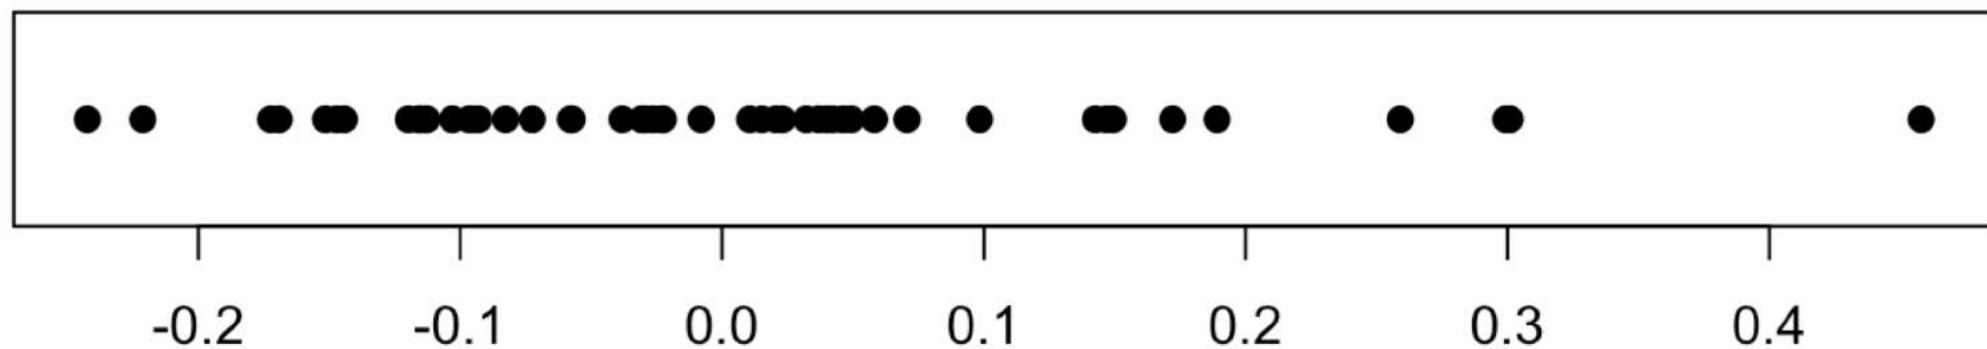
$$\begin{matrix} U \\ 2 \times 2 \end{matrix} \times \begin{matrix} \text{Sigma} \\ 2 \times 2 \end{matrix} \times \begin{matrix} V^T \\ 2 \times 49 \end{matrix}$$

$$\begin{matrix} U \\ 2 \times 2 \end{matrix} \times \begin{matrix} \text{Sigma} \\ 2 \times 1 \end{matrix} \times \begin{matrix} V^T \\ 1 \times 49 \end{matrix}$$

# One Dimensional Song Data



# One Dimensional Song Data



# Projection

We use the following formula to project our data onto the ellipse axes:

$$P = U\Sigma X$$



```
training_projections = svd$u %*% diag(svd$d) %*% centered
```

# Projections



```
[,1]      [,2]      [,3]      [,4]
energy  0.2195102 0.1435102 -0.0964898 -0.0834898
loudness 0.2020816 1.7200816 -1.9869184 -0.5579184
[,1]      [,2]      [,3]      [,4]
PC1 -0.03816334 -0.1690882 0.1890630 0.05818034
PC2 -0.31990593 -0.2002474 0.1299837 0.11897906
```

# Does this song belong?



# Step 1: Center and Project the Test Song



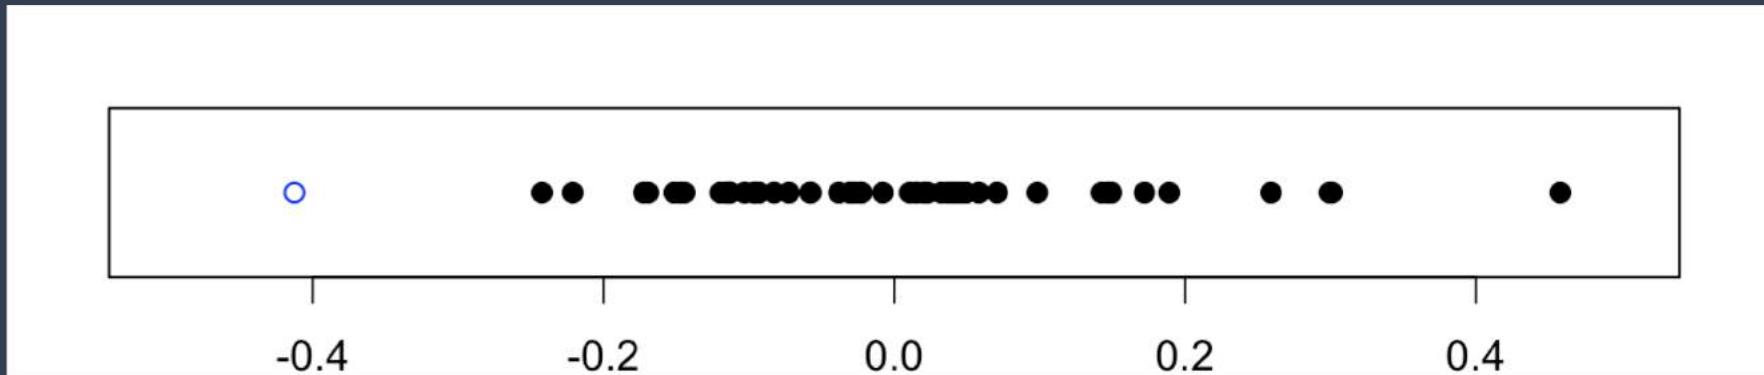
```
test_song = c(-2, 2) - rowMeans(training_set)
test_projections = svd$u %*% diag(svd$d) %*% test_song
```

# Step 2: Calculate Distances



```
distancess = abs(test_projections[1,] - training_projections[1,])
dist = data.frame(
  "min.dist" = min(distancess),
  "min.index" = which(distancess == min(distancess)))
)
```

# Step 3: Interpret Results



# Reminder

The value of PCA is in its ability to reduce the number of dimensions we need to work with.

# All Song Features

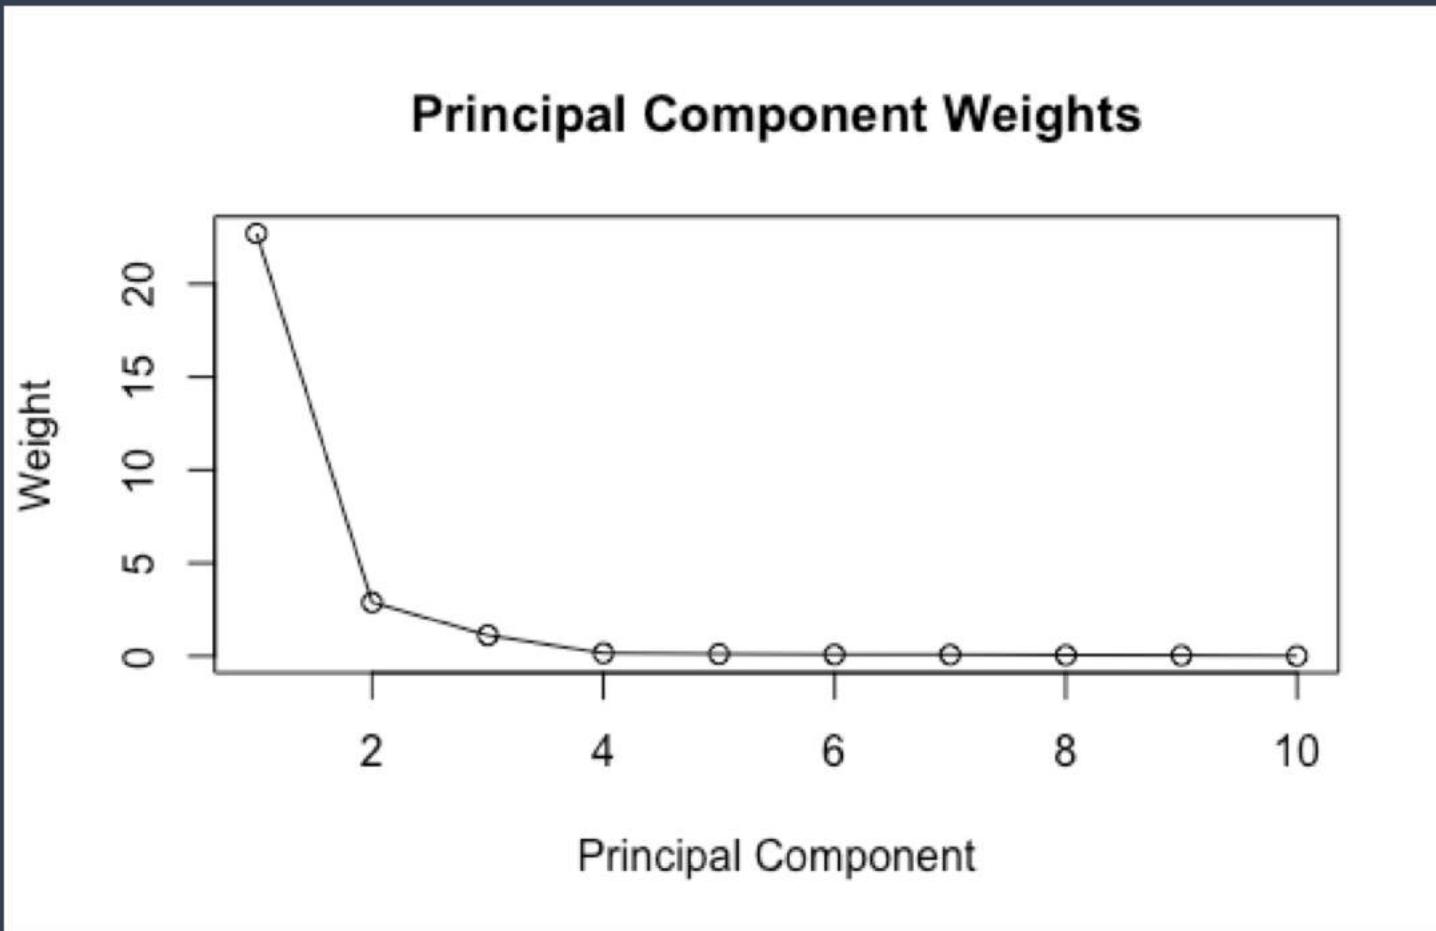


```
songs = read.csv("songdata.csv")
matrix = t(data.matrix(songs)[,-1])

centered = matrix - rowMeans(matrix)

svd = svd(centered/9)
```

# All Song Features

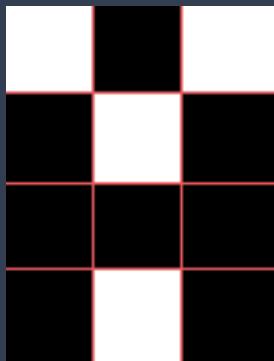


# Face Dimensional Reduction

What if the number of dimensions gets even bigger?

# Images as Vectors

# Representing Images as Vectors



This 4 by 3 image becomes a 12 by 1 image.

Each pixel represents a feature.



# Face Images



Each image is 342 pixels by 305 pixels.

Each image has been normalized to the best of my abilities.

This will produce a vector that is 104310 pixels long.

That's 104310 features to consider.

# The Training Set



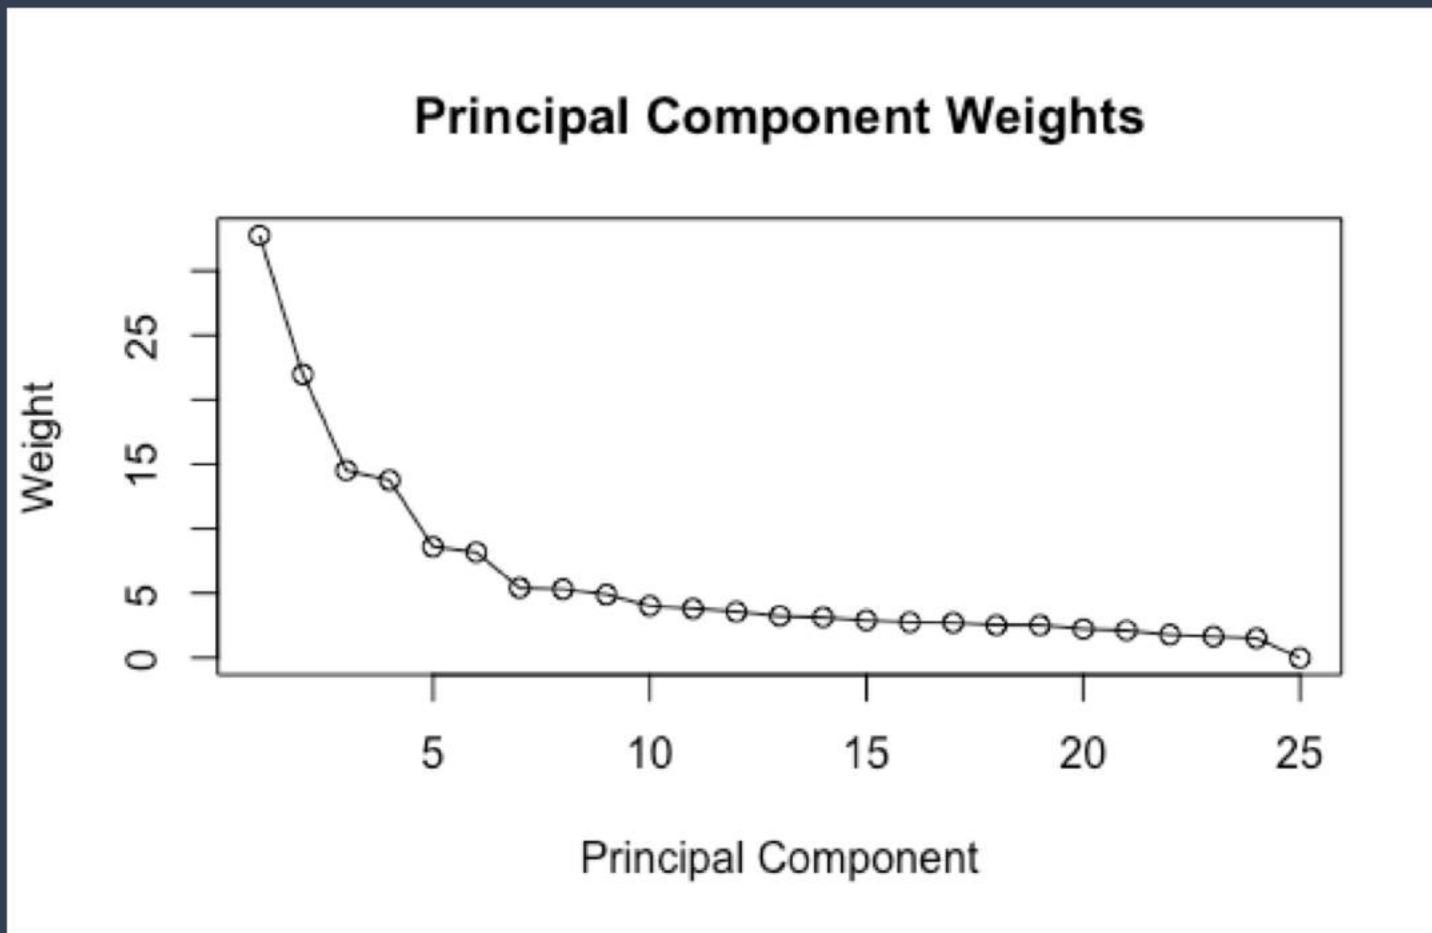
# Normalize the Training Set



# Normalize the Training Set



# Principal Component Weights



# Principal Component Vectors



# What does that mean?



# The Projections



# Let's Test the Model



Image In Set



Person In Set



Not In Set

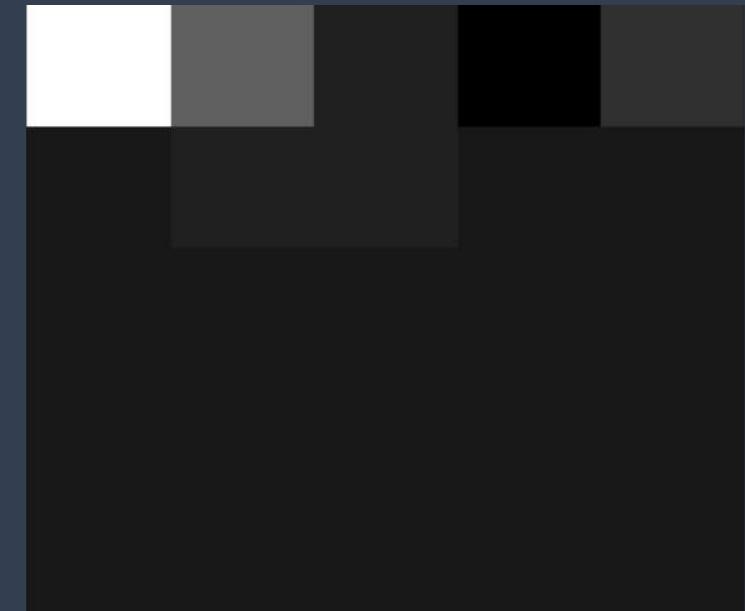
# Projections



**Image In Set**



**Person In Set**



**Not In Set**

# Side by Side Comparison



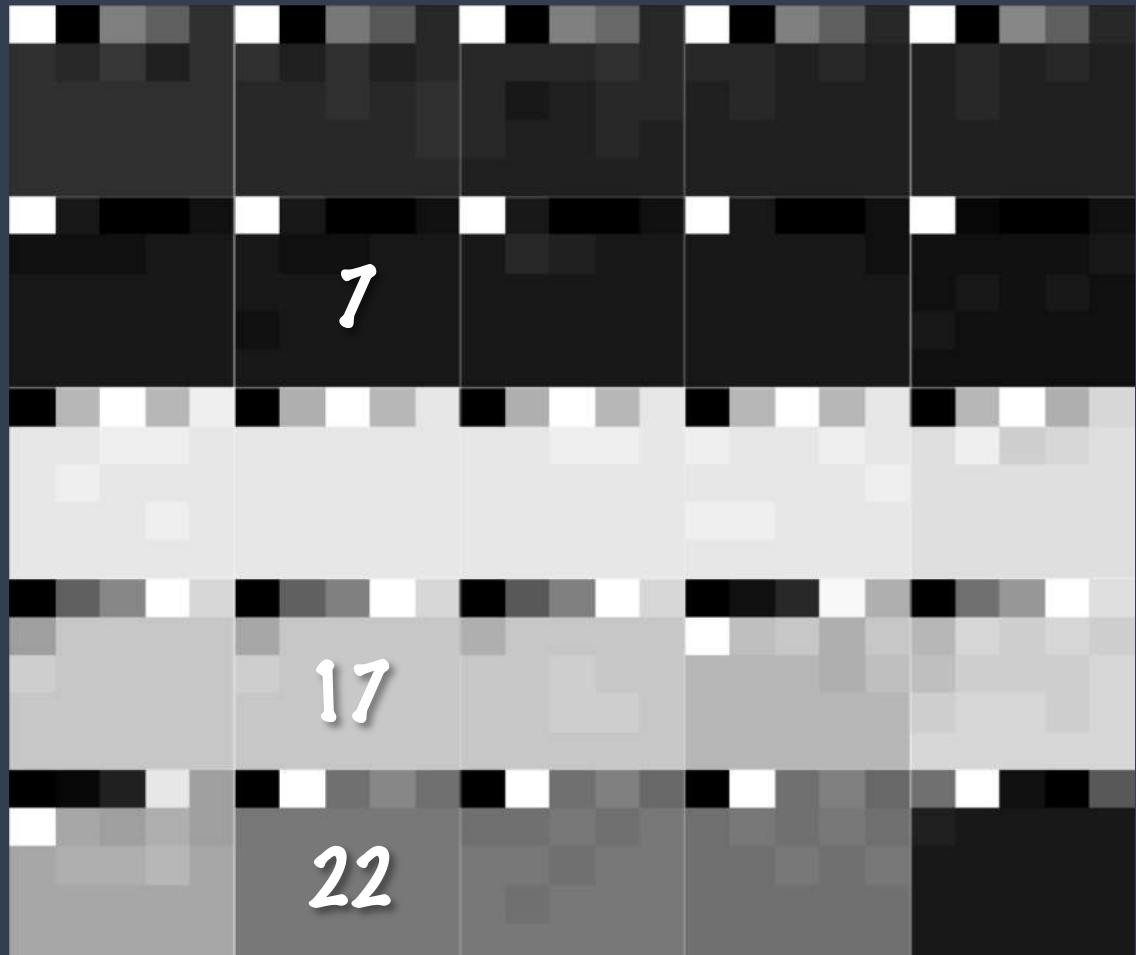
**Image In Set**  
**Closest to Image 17**  
**Distance: 0**



**Person In Set**  
**Closest to Image 22**  
**Distance: 35.98396**



**Person Not In Set**  
**Closest to Image 7**  
**Distance: 527.8126**



# Side by Side Comparison



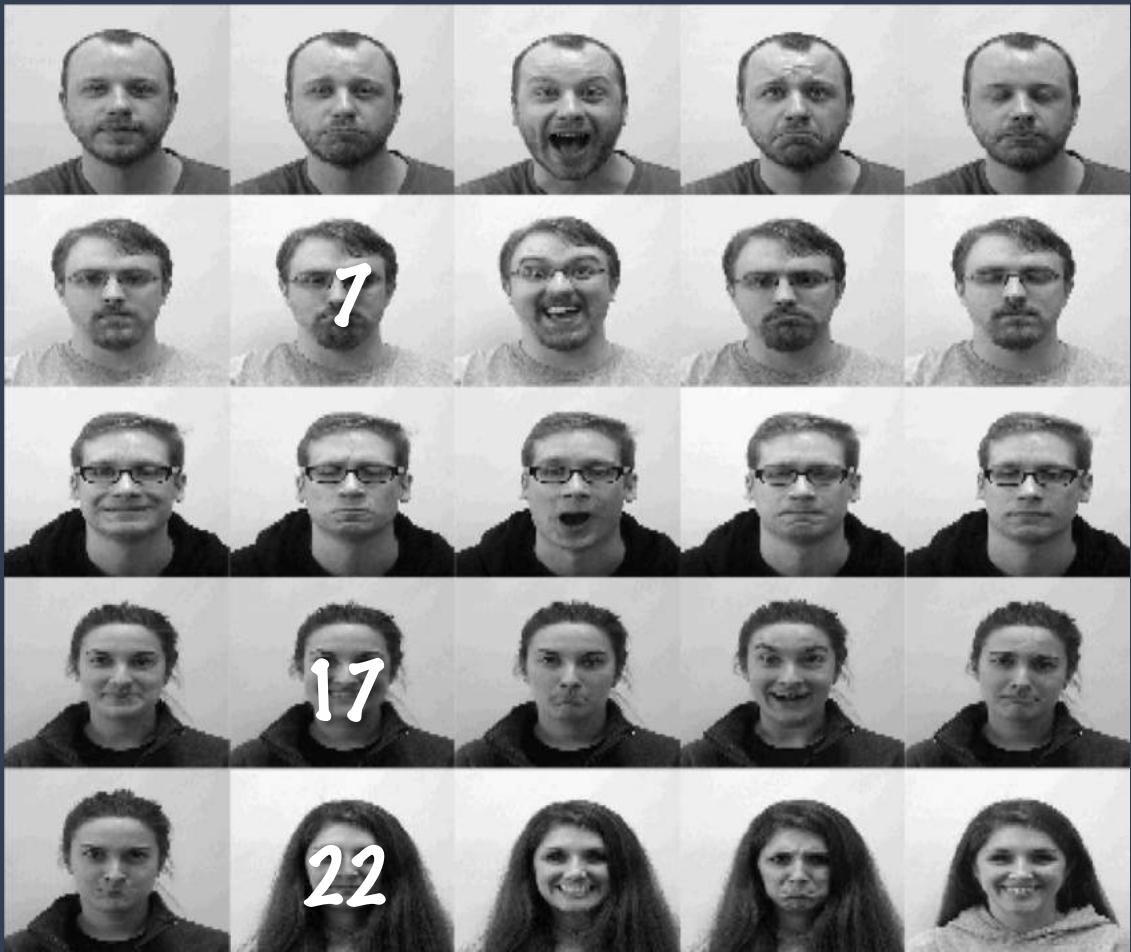
**Image In Set**  
**Closest to Image 17**  
**Distance: 0**



**Person In Set**  
**Closest to Image 22**  
**Distance: 35.98396**



**Person Not In Set**  
**Closest to Image 7**  
**Distance: 527.8126**



# Here's the Process

# Step 1: Build the Training Set



```
training_set_files = list.files(path = 'training_set', pattern = '.jpg', full.names = TRUE)
training_set_faces = map_il(training_set_files, load.image)
training_set = matrix(
  unlist(lapply(training_set_faces, function(x) x[, , , 1])),
  ncol = length(training_set_faces),
  byrow = FALSE)
```

# Step 2: Mean Normalization



```
average_face = rowMeans(training_set)  
centered = training_set - average_face
```

# Step 3: Singular Value Decomposition



```
svd = svd(centered / sqrt(ncol(training_set))))
```

# Step 4: Calculate the Projections



```
training_projections = t(svd$u %*% diag(svd$d)) %*% centered
```

# All Together Now

```
● ● ●

# Step 1 Create Training Set
training_set_files = list.files(path = 'training_set', pattern = '.jpg', full.names = TRUE)
training_set_faces = map_l1(training_set_files, load.image)
training_set = matrix(
  unlist(lapply(training_set_faces, function(x) x[, , , 1])),
  ncol = length(training_set_faces),
  byrow = FALSE)

# Step 2 Perform Mean Normalization
average_face = rowMeans(training_set)
centered = training_set - average_face

# Step 3 Perform SVD
svd = svd(centered / sqrt(ncol(training_set)))

# Step 4 Project the Training Set Into Low-Rank Dimension
training_projections = t(svd$u %*% diag(svd$d)) %*% centered
```

# Test The Model

# Step 1: Create Test Set



```
test_files = list.files(path = 'test_set', pattern = '.jpg', full.names = TRUE)
test_faces = map_il(test_files, load.image)
test_set = matrix(
  unlist(lapply(test_faces, function(x) x[, , , 1] - average_face)),
  ncol = 3,
  byrow = FALSE)
```

# Step 2: Project the Test Set



```
test_set_projections = t(svd$u %*% diag(svd$d)) %*% test_set
```

# Step 3: Calculate the Distances



```
distances = apply(test_set_projections, 2, function(test) {  
  dist = apply(training_projections, 2, function(x) sqrt(sum((x - test) ^ 2)))  
  data.frame(  
    "min.index" = which(dist == min(dist)),  
    "min.dist" = min(dist),  
    "max.index" = which(dist == max(dist)),  
    "max.dist" = max(dist)  
  )  
})
```

# Step 4: See the Results

```
● ● ●

[[1]]
  min.index min.dist max.index max.dist
1          17           0          8 2373.384

[[2]]
  min.index min.dist max.index max.dist
1          22 35.98396           6 2734.868

[[3]]
  min.index min.dist max.index max.dist
1          7 527.8126          12 2857.625
```

# We Have Problems



# Steve Crow

Java Developer Advocate  
Nexmo



@cr0wst



# Questions?

# Links and Stuff

<https://towardsdatascience.com/a-one-stop-shop-for-principal-component-analysis-5582fb7e0a9c>

<https://machinelearningmastery.com/singular-value-decomposition-for-machine-learning/>

<https://github.com/cr0wst/making-faces>